

# Ninjacart Category Agent System

## Complete Technical Implementation Guide

Version: 1.0

Document Type: Technical Specification & Implementation Guide

Last Updated: October 2025

Target Audience: Product Managers, Data Scientists, Backend Engineers, Frontend Engineers, DevOps Engineers

**Document Purpose:** This comprehensive guide provides complete technical specifications, implementation details, code examples, and deployment instructions for building Ninjacart's AI-powered Category Agent System.

## Table of Contents

- [1. Executive Summary](#)
- [2. System Architecture Overview](#)
- [3. Technology Stack](#)
- [4. Infrastructure Setup](#)
- [5. Data Layer Implementation](#)
- [6. Memory Systems Implementation](#)
- [7. Agent Implementation - Complete](#)
- [8. Orchestrator Implementation](#)
- [9. API Gateway & Backend Services](#)
- [10. User Interface Development](#)
- [11. Testing Strategy](#)
- [12. Deployment & Operations](#)
- [13. Monitoring & Observability](#)
- [14. Team Responsibilities](#)
- [15. Implementation Timeline](#)
- [16. Appendices](#)

## 1. Executive Summary

### 1.1 System Overview

The Ninjacart Category Agent System is an AI-powered platform that automates category management decisions including:

- Strategic Planning:** Weekly analysis and target setting (Every Saturday)
- Daily Pricing:** Automated SP recommendations with explainable reasoning (Every morning 7 AM)
- Demand Forecasting:** ML-based predictions (85-90% accuracy)
- Campaign Management:** Intelligent segment targeting and ROI optimization
- Real-time Monitoring:** Proactive interventions and alerts

## 1.2 Key Metrics & Goals

Metric	Current (Manual)	Target (Automated)	Impact
Pricing Decision Time	30-45 min/SKU	2 min/SKU	93% reduction
Forecast Accuracy	~70%	85-90%	+15-20pp improvement
Campaign Success Rate	60-65%	75-80%	+15pp improvement
Manual Effort	4-5 hours/day	1 hour/day	75% reduction
Revenue Protection	Baseline	+₹10-15L/day	Incremental

## 1.3 Architecture Layers



USER LAYER (Chat, Dashboard, Approval Queue)



API GATEWAY (FastAPI, Auth, Rate Limiting)



ORCHESTRATION (LangGraph, Prefect)



AGENT LAYER (8 Specialized Agents)

Strategic: Weekly Planning

Tactical: Demand, Inventory, Pricing, Campaign,

Monitoring, Returns



MEMORY & DATA (Redis, PostgreSQL, Pinecone)

## 2. System Architecture Overview

### 2.1 High-Level Architecture Diagram



## USER LAYER

Chat Interface | | Poker Dashboard | | Approval Queue |  
(Streamlit) | | (React) | | (React) | |



## API GATEWAY LAYER

FastAPI Gateway (Authentication, Rate Limiting, Routing) | |  
- JWT Authentication | |  
- Request/Response logging | | |  
- Error handling & retries | |



## ORCHESTRATION LAYER

Orchestrator Agent (LangGraph / Prefect) | |  
- Intent detection (NLP) | |  
- Agent routing & coordination | | |  
- Workflow execution (DAG-based) | | |  
- Error handling & retries | |  
- Human-in-loop approval management | |

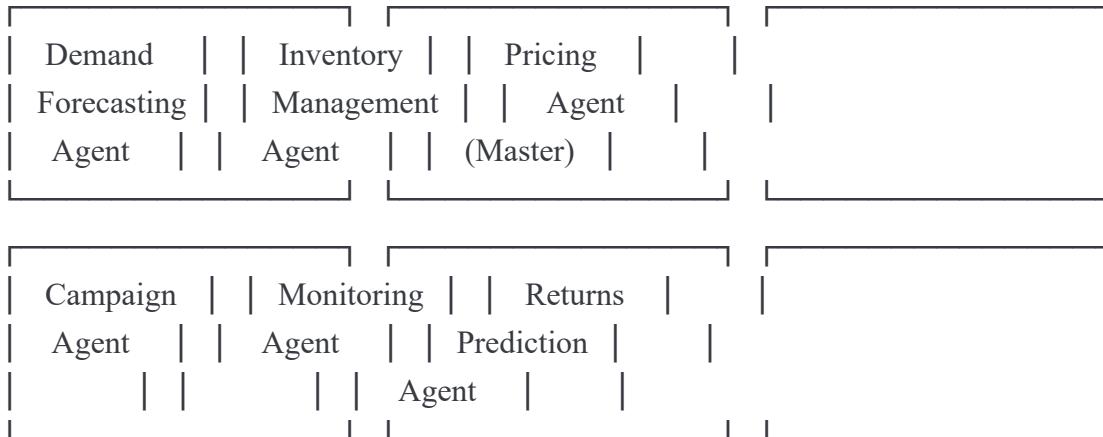


## AGENT LAYER

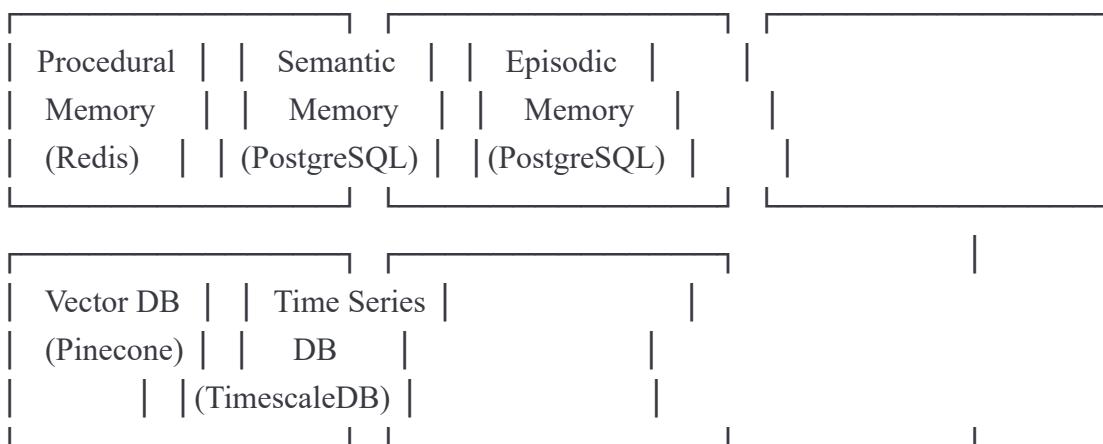
### STRATEGIC AGENT (Weekly):

Weekly Planning Agent | |  
- Analyze past week performance | | |  
- Assess supply stability | |  
- Decide strategy (GROW\_TPC vs STABILIZE) | | |  
- Set weekly targets | |  
- Plan campaigns | |

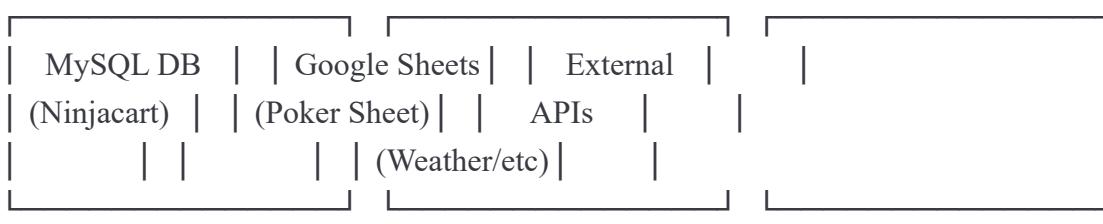
## TACTICAL AGENTS (Daily):



## MEMORY & DATA LAYER



## DATA SOURCES LAYER



## 2.2 Data Flow Patterns

### Request-Response Flow (User-Initiated):



User Request → API Gateway → Orchestrator →  
Agent Selection → Agent Execution → Memory Access →  
Result Compilation → Response to User

### Scheduled Flow (Automated):



Cron Trigger → Orchestrator → Agent Execution →  
Memory Update → Notification System

### Learning Loop:



Agent Decision → Execution → Outcome Measurement →  
Episodic Memory → Vector DB → Future RAG Retrieval

---

# 3. Technology Stack

## 3.1 Core Technologies

Component	Technology	Version	Justification
Primary Language	Python	3.11+	Mature ML ecosystem, async support, type hints
LLM	OpenAI GPT-4 / Claude Sonnet	Latest	SOTA reasoning, low hallucination rate
Agent Framework	LangChain + LangGraph	0.1.0 / 0.0.20	Graph-based workflows, observability
Orchestration	Prefect	2.14.0	Workflow management, scheduling, monitoring
Backend API	FastAPI	0.104.1	High performance, async, auto OpenAPI docs
Cache/Config	Redis	7.x	In-memory speed, pub/sub, TTL support
Database	PostgreSQL	15.x	ACID compliance, JSONB, mature
Time Series	TimescaleDB	Latest	Optimized time-series queries
Vector DB	Pinecone	Latest	Managed, scalable, low latency
ML Framework	Scikit-learn, XGBoost, Prophet	Latest	Industry standard, proven accuracy
Frontend	React + Streamlit	18.x / Latest	Complex UI (React), Rapid prototyping (Streamlit)
Monitoring	Prometheus + Grafana	Latest	Open-source, extensive integrations
Logging	ELK Stack	Latest	Centralized logging, search, analysis

## 3.2 Python Dependencies

requirements.txt:



```
# LLM & Agent Framework
```

```
langchain==0.1.0
```

```
langgraph==0.0.20
```

```
openai==1.6.1
```

```
anthropic==0.8.1
```

```
# Orchestration & Workflow
```

```
prefect==2.14.0
```

```
celery==5.3.4
```

```
redis==5.0.1
```

```
# API Framework
```

```
fastapi==0.104.1
```

```
uvicorn[standard]==0.24.0
```

```
pydantic==2.5.0
```

```
python-multipart==0.0.6
```

```
# Database & ORM
```

```
sqlalchemy==2.0.23
```

```
alembic==1.13.0
```

```
psycopg2-binary==2.9.9
```

```
asyncpg==0.29.0
```

```
# Vector Database
```

```
pinecone-client==2.2.4
```

```
# ML & Data Science
```

```
scikit-learn==1.3.2
```

```
xgboost==2.0.2
```

```
prophet==1.1.5
```

```
pandas==2.1.4
```

```
numpy==1.26.2
```

```
scipy==1.11.4
```

```
# Google Sheets Integration
```

```
google-api-python-client==2.111.0
```

```
google-auth==2.25.2
```

```
google-auth-oauthlib==1.2.0
```

```
google-auth-httplib2==0.2.0
```

```
# HTTP & API Clients
```

```
httpx==0.25.2
requests==2.31.0

# Utilities
python-dotenv==1.0.0
tenacity==8.2.3 # Retry logic
pyyaml==6.0.1
python-json-logger==2.0.7

# Monitoring & Metrics
prometheus-client==0.19.0
opentelemetry-api==1.21.0
opentelemetry-sdk==1.21.0
```

```
# Testing
pytest==7.4.3
pytest-asyncio==0.21.1
pytest-cov==4.1.0
httpx==0.25.2 # For testing FastAPI
```

```
# Security
python-jose[cryptography]==3.3.0
passlib[bcrypt]==1.7.4
```

---

## 4. Infrastructure Setup

### 4.1 Development Environment

**Prerequisites:**

- Python 3.11+
- Docker & Docker Compose
- Git
- 16GB RAM minimum
- (Optional) GPU for local LLM testing

### 4.2 Docker Compose Setup

**docker-compose.yml:**



version: '3.8'

services:

# PostgreSQL with TimescaleDB

postgres:

image: timescale/timescaledb:latest-pg15

container\_name: category\_agent\_db

environment:

POSTGRES\_DB: category\_agent

POSTGRES\_USER: admin

POSTGRES\_PASSWORD: \${DB\_PASSWORD}

POSTGRES\_INITDB\_ARGS: "-c shared\_preload\_libraries=timescaledb"

ports:

- "5432:5432"

volumes:

- postgres\_data:/var/lib/postgresql/data

- ./init\_scripts:/docker-entrypoint-initdb.d

healthcheck:

test: ["CMD-SHELL", "pg\_isready -U admin -d category\_agent"]

interval: 10s

timeout: 5s

retries: 5

# Redis

redis:

image: redis:7-alpine

container\_name: category\_agent\_redis

ports:

- "6379:6379"

volumes:

- redis\_data:/data

command: redis-server --appendonly yes

healthcheck:

test: ["CMD", "redis-cli", "ping"]

interval: 10s

timeout: 3s

retries: 5

# FastAPI Backend

api:

build:

```
context: ./backend
dockerfile: Dockerfile
container_name: category_agent_api
ports:
- "8000:8000"
environment:
- DATABASE_URL=postgresql+asyncpg://admin:${DB_PASSWORD}@postgres:5432/category_agent
- REDIS_URL=redis://redis:6379
- OPENAI_API_KEY=${OPENAI_API_KEY}
- PINECONE_API_KEY=${PINECONE_API_KEY}
- PINECONE_ENV=${PINECONE_ENV}
depends_on:
  postgres:
    condition: service_healthy
  redis:
    condition: service_healthy
volumes:
- ./backend:/app
- ./models:/app/models
command: uvicorn main:app --host 0.0.0.0 --port 8000 --reload
```

#### # Streamlit Frontend

```
streamlit:
build:
  context: ./frontend
  dockerfile: Dockerfile.streamlit
container_name: category_agent_streamlit
ports:
- "8501:8501"
environment:
- API_URL=http://api:8000
depends_on:
- api
volumes:
- ./frontend/streamlit_app:/app
```

#### # React Dashboard

```
react_dashboard:
build:
  context: ./frontend/react_dashboard
  dockerfile: Dockerfile
```

```
container_name: category_agent_dashboard
ports:
- "3000:3000"
environment:
- REACT_APP_API_URL=http://localhost:8000
volumes:
- ./frontend/react_dashboard:/app
- /app/node_modules
```

# Prefect Server (Workflow Orchestration)

```
prefect_server:
image: prefecthq/prefect:2.14.0-python3.11
container_name: category_agent_prefect
ports:
- "4200:4200"
environment:
- PREFECT_API_URL=http://0.0.0.0:4200/api
- PREFECT_SERVER_API_HOST=0.0.0.0
command: prefect server start
volumes:
- prefect_data:/root/.prefect
```

# Prometheus (Metrics)

```
prometheus:
image: prom/prometheus:latest
container_name: category_agent_prometheus
ports:
- "9090:9090"
volumes:
- ./monitoring/prometheus.yml:/etc/prometheus/prometheus.yml
- prometheus_data:/prometheus
command:
- '--config.file=/etc/prometheus/prometheus.yml'
- '--storage.tsdb.path=/prometheus'
```

# Grafana (Visualization)

```
grafana:
image: grafana/grafana:latest
container_name: category_agent_grafana
ports:
- "3001:3000"
```

**environment:**

- GF\_SECURITY\_ADMIN\_PASSWORD=\${GRAFANA\_PASSWORD}

**volumes:**

- grafana\_data:/var/lib/grafana

- ./monitoring/grafana/dashboards:/etc/grafana/provisioning/dashboards

**depends\_on:**

- prometheus

**volumes:**

- postgres\_data:

- redis\_data:

- prefect\_data:

- prometheus\_data:

- grafana\_data:

## 4.3 Environment Configuration

**.env:**



bash

```
# Database
DB_PASSWORD=your_secure_password_here
DATABASE_URL=postgresql+asyncpg://admin:${DB_PASSWORD}@localhost:5432/category_agent

# Redis
REDIS_URL=redis://localhost:6379

# LLM APIs
OPENAI_API_KEY=sk-...
ANTHROPIC_API_KEY=sk-ant-...

# Vector DB
PINECONE_API_KEY=...
PINECONE_ENV=us-west1-gcp
PINECONE_INDEX=category-agent-episodes

# Ninjacart Data Sources
NINJACART_DB_HOST=your-mysql-host.rds.amazonaws.com
NINJACART_DB_PORT=3306
NINJACART_DB_USER=readonly_user
NINJACART_DB_PASSWORD=...
NINJACART_DB_NAME=ninjacart_production

# Google Sheets
GOOGLE_SHEETS_CREDENTIALS_PATH=/app/credentials/google_credentials.json
POKER_SHEET_ID=130Q4yw1nfwk9Mj839NbAdTb3dtdCXVRxf0WNpYhYQd8

# External APIs
WEATHER_API_KEY=...
WEATHER_API_URL=https://api.openweathermap.org/data/2.5

# JWT Authentication
JWT_SECRET_KEY=your_jwt_secret_key_here
JWT_ALGORITHM=HS256
JWT_ACCESS_TOKEN_EXPIRE_MINUTES=1440

# Monitoring
PROMETHEUS_PORT=9090
GRAFANA_PORT=3001
GRAFANA_PASSWORD=admin
```

```
# Application  
APP_ENV=development # development, staging, production  
LOG_LEVEL=INFO
```

## 4.4 Initialization Scripts

**init\_scripts/01\_create\_extensions.sql:**



```
-- Enable TimescaleDB extension  
CREATE EXTENSION IF NOT EXISTS timescaledb;  
  
-- Enable UUID generation  
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";  
  
-- Enable pg_stat_statements for query performance monitoring  
CREATE EXTENSION IF NOT EXISTS pg_stat_statements;
```

---

## 5. Data Layer Implementation

### 5.1 Complete Database Schema

**schema.sql:**



```
-- =====
-- SEMANTIC MEMORY: Domain Knowledge & SKU Characteristics
-- =====

CREATE TABLE sku_knowledge (
    sku_id VARCHAR(50) PRIMARY KEY,
    sku_name VARCHAR(255) NOT NULL,
    category VARCHAR(100) NOT NULL,

    -- Target metrics
    daily_penetration_target_min DECIMAL(5,2),
    daily_penetration_target_max DECIMAL(5,2),
    weekly_penetration_target DECIMAL(5,2),
    gm_target_min DECIMAL(8,2),
    gm_target_max DECIMAL(8,2),
    tpo_target DECIMAL(8,2),

    -- SKU Characteristics
    crop_cycle_weeks INTEGER,
    shelf_life_days INTEGER,
    peak_season VARCHAR[] DEFAULT '{}',
    price_elasticity DECIMAL(4,2),
    festival_lift_pct DECIMAL(5,2),
    rain_impact_pct DECIMAL(5,2),
    quality_sensitive BOOLEAN DEFAULT TRUE,

    -- Formulas (JSON for flexibility)
    pricing_formulas JSONB DEFAULT '{}',
    margin_rules JSONB DEFAULT '{}',

    -- Metadata
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);


```

```
CREATE INDEX idx_sku_category ON sku_knowledge(category);
```

```
-- =====
-- EPISODIC MEMORY: Historical Decisions & Outcomes
-- =====
```

```
CREATE TABLE pricing_episodes (
    episode_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    sku_id VARCHAR(50) REFERENCES sku_knowledge(sku_id),
    episode_date DATE NOT NULL,
    -- Context
    wsp DECIMAL(8,2),
    pp DECIMAL(8,2),
    inventory_kg DECIMAL(10,2),
    inventory_surplus_pct DECIMAL(5,2),
    pm DECIMAL(8,2),
    supply_status VARCHAR(20), -- SURPLUS, BALANCED, DEFICIT
    day_of_week INTEGER,
    is_festival BOOLEAN,
    weather_condition VARCHAR(50),
    -- Decision
    sp_recommended DECIMAL(8,2),
    sp_approved DECIMAL(8,2),
    sp_reasoning TEXT,
    confidence DECIMAL(4,3),
    human_approved BOOLEAN,
    approved_by VARCHAR(100),
    approval_timestamp TIMESTAMP,
    -- Expected Outcomes
    tpc_predicted DECIMAL(10,2),
    penetration_predicted DECIMAL(5,2),
    gm_predicted DECIMAL(8,2),
    -- Actual Outcomes (filled after day completes)
    tpc_actual DECIMAL(10,2),
    penetration_actual DECIMAL(5,2),
    gm_actual DECIMAL(8,2),
    orders_actual INTEGER,
    -- Learning Metrics
    tpc_error_pct DECIMAL(5,2),
    penetration_error_pp DECIMAL(5,2),
    gm_error_pct DECIMAL(5,2),
    overall_accuracy DECIMAL(4,3),
```

```
learning_notes TEXT,
```

```
-- Metadata
```

```
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
```

```
);
```

```
CREATE INDEX idx_episode_sku_date ON pricing_episodes(sku_id, episode_date);
```

```
CREATE INDEX idx_episode_date ON pricing_episodes(episode_date);
```

```
CREATE INDEX idx_episode_supply_status ON pricing_episodes(supply_status);
```

```
-- ======  
-- WEEKLY PLANS: Strategic Direction  
-- ======
```

```
CREATE TABLE weekly_plans (
```

```
plan_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
sku_id VARCHAR(50) REFERENCES sku_knowledge(sku_id),  
week_start_date DATE NOT NULL,  
week_end_date DATE NOT NULL,  
week_number INTEGER, -- ISO week number
```

```
-- Previous Week Performance
```

```
prev_week_tpc_actual DECIMAL(10,2),  
prev_week_tpc_target DECIMAL(10,2),  
prev_week_achievement_pct DECIMAL(5,2),  
prev_week_penetration_actual DECIMAL(5,2),  
prev_week_penetration_target DECIMAL(5,2),  
prev_week_gm_avg DECIMAL(8,2),
```

```
-- Supply Assessment
```

```
supply_stability VARCHAR(20), -- STABLE, UNSTABLE  
wsp_cv_pct DECIMAL(5,2),  
wsp_mean DECIMAL(8,2),  
wsp_range_min DECIMAL(8,2),  
wsp_range_max DECIMAL(8,2),  
quality_score DECIMAL(4,3),  
qqc_rejection_avg_pct DECIMAL(5,2),  
vendor_pipeline_count INTEGER,  
vendor_reliability_score DECIMAL(4,3),  
next_week_risk_score DECIMAL(4,3),
```

next\_week\_risk\_level VARCHAR(20), -- LOW, MODERATE, HIGH

-- Strategic Decision

strategy VARCHAR(50), -- GROW\_TPC, STABILIZE\_TPC\_IMPROVE\_MARGINS  
strategy\_focus VARCHAR(50), -- VOLUME\_GROWTH, MARGIN\_PRESERVATION  
strategy\_rationale TEXT,  
margin\_approach VARCHAR(50),

-- Historical Benchmark

similar\_weeks\_count INTEGER,  
mode\_tpc\_kg\_per\_day DECIMAL(10,2),  
median\_tpc\_kg\_per\_day DECIMAL(10,2),  
p75\_tpc\_kg\_per\_day DECIMAL(10,2),  
safe\_baseline\_kg\_per\_day DECIMAL(10,2),

-- Next Week Targets

daily\_tpc\_target\_kg DECIMAL(10,2),  
weekly\_tpc\_target\_kg DECIMAL(10,2),  
growth\_vs\_baseline\_pct DECIMAL(5,2),  
penetration\_target\_pct DECIMAL(5,2),  
tpo\_target\_kg DECIMAL(8,2),  
gm\_target\_min DECIMAL(8,2),  
gm\_target\_max DECIMAL(8,2),  
gm\_preference VARCHAR(50),  
pm\_target\_min DECIMAL(8,2),  
pm\_target\_max DECIMAL(8,2),  
sm\_target\_min DECIMAL(8,2),  
sm\_target\_max DECIMAL(8,2),

-- Penetration Gap Analysis

ideal\_penetration\_pct DECIMAL(5,2),  
current\_penetration\_pct DECIMAL(5,2),  
penetration\_gap\_pp DECIMAL(5,2),  
campaign\_required BOOLEAN,

-- Planned Campaign (JSONB for flexibility)

planned\_campaign JSONB,

-- Success Checkpoints

checkpoints JSONB,

```
-- Approval
approval_status VARCHAR(20), -- PENDING, APPROVED, REJECTED, MODIFIED
approved_by VARCHAR(100),
approved_at TIMESTAMP,
rejection_reason TEXT,

-- Metadata
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
CREATE INDEX idx_weekly_plan_sku_week ON weekly_plans(sku_id, week_start_date);
CREATE INDEX idx_weekly_plan_week ON weekly_plans(week_start_date);
CREATE INDEX idx_weekly_plan_approval ON weekly_plans(approval_status);
```

```
-- =====
-- CAMPAIGN HISTORY: Campaign Performance
-- =====
```

```
CREATE TABLE campaign_episodes (
    campaign_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    sku_id VARCHAR(50) REFERENCES sku_knowledge(sku_id),
    campaign_date DATE NOT NULL,
    campaign_type VARCHAR(50), -- SKU_PROMOTION, CART_DISCOUNT, MILESTONE, SLAB_PRICING
    campaign_trigger VARCHAR(50), -- PENETRATION_DROP, EXCESS_INVENTORY, PLANNED_WEEKLY
```

```
-- Segment Details
segment_name VARCHAR(100),
segment_size INTEGER,
segment_definition JSONB,
```

```
-- Discount Structure
discount_type VARCHAR(50), -- FLAT, PERCENTAGE, SLAB
discount_structure JSONB,
```

```
-- Timing
start_timestamp TIMESTAMP,
end_timestamp TIMESTAMP,
duration_hours INTEGER,
```

```
-- Predictions
```

```
conversion_rate_predicted DECIMAL(5,4),
customers_converting_predicted INTEGER,
avg_order_size_predicted DECIMAL(8,2),
incremental_tpc_predicted DECIMAL(10,2),
penetration_lift_predicted_pp DECIMAL(5,2),
budget_allocated DECIMAL(10,2),
expected_roi DECIMAL(8,2),
```

-- *Actuals*

```
conversion_rate_actual DECIMAL(5,4),
customers_converting_actual INTEGER,
avg_order_size_actual DECIMAL(8,2),
incremental_tpc_actual DECIMAL(10,2),
penetration_lift_actual_pp DECIMAL(5,2),
discount_burn_actual DECIMAL(10,2),
incremental_revenue_actual DECIMAL(10,2),
actual_roi DECIMAL(8,2),
```

-- *Performance Metrics*

```
conversion_prediction_accuracy DECIMAL(5,2),
tpc_prediction_accuracy DECIMAL(5,2),
roi_prediction_accuracy DECIMAL(5,2),
overall_success_score DECIMAL(4,3),
```

-- *Learning*

```
learning_notes TEXT,
```

-- *Metadata*

```
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
CREATE INDEX idx_campaign_sku_date ON campaign_episodes(sku_id, campaign_date);
CREATE INDEX idx_campaign_type ON campaign_episodes(campaign_type);
CREATE INDEX idx_campaign_trigger ON campaign_episodes(campaign_trigger);
```

```
-- =====
-- TIME SERIES DATA: Daily Actuals (TimescaleDB Hypertable)
-- =====
```

```
CREATE TABLE daily_actuals (
```

```
ts TIMESTAMPTZ NOT NULL,  
sku_id VARCHAR(50) NOT NULL,
```

-- Sales Metrics

```
orders_count INTEGER,  
customers_count INTEGER,  
tpc DECIMAL(10,2),  
tpo DECIMAL(8,2),  
penetration DECIMAL(5,2),
```

-- Pricing

```
wsp DECIMAL(8,2),  
pp DECIMAL(8,2),  
sp DECIMAL(8,2),
```

-- Margins

```
pm DECIMAL(8,2),  
sm DECIMAL(8,2),  
cm DECIMAL(8,2),  
gm DECIMAL(8,2),
```

-- Inventory

```
inventory_kg DECIMAL(10,2),  
supply_status VARCHAR(20),  
days_of_inventory DECIMAL(4,2),
```

-- Quality

```
qqc_rejection_pct DECIMAL(5,2),  
otpr_pct DECIMAL(5,2),
```

-- External Factors

```
is_festival BOOLEAN,  
is_weekend BOOLEAN,  
weather_condition VARCHAR(50),  
temperature DECIMAL(4,1),
```

PRIMARY KEY (ts, sku\_id)

);

-- Convert to hypertable (TimescaleDB specific)  
SELECT create\_hypertable('daily\_actuals', 'ts',

```

chunk_time_interval => INTERVAL '7 days',
if_not_exists => TRUE
);

-- Create indexes
CREATE INDEX idx_daily_sku ON daily_actuals(sku_id, ts DESC);
CREATE INDEX idx_daily_date ON daily_actuals(ts DESC);

-- =====
-- DEMAND FORECASTS
-- =====

CREATE TABLE demand_forecasts (
    forecast_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    sku_id VARCHAR(50) NOT NULL,
    forecast_date DATE NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

    -- Model Metadata
    model_version VARCHAR(50),
    model_type VARCHAR(50), -- XGBOOST, PROPHET, ENSEMBLE

    -- Base Forecast
    demand_forecast_kg DECIMAL(10,2),
    confidence DECIMAL(4,3),

    -- Uncertainty Bands
    p50_kg DECIMAL(10,2),
    p75_kg DECIMAL(10,2),
    p90_kg DECIMAL(10,2),

    -- Segment Breakdown
    retail_forecast_kg DECIMAL(10,2),
    bulk_forecast_kg DECIMAL(10,2),

    -- Features Used (for debugging)
    features_snapshot JSONB,

    -- Adjustments Applied
    adjustments_applied JSONB,

```

```
-- Actuals (filled post-day)
actual_demand_kg DECIMAL(10,2),
forecast_error_pct DECIMAL(5,2),
forecast_error_kg DECIMAL(10,2),

UNIQUE (sku_id, forecast_date, created_at)
);
```

```
CREATE INDEX idx_forecast_sku_date ON demand_forecasts(sku_id, forecast_date);
CREATE INDEX idx_forecast_date ON demand_forecasts(forecast_date);
```

```
-- =====
-- AGENT EXECUTION LOGS
-- =====
```

```
CREATE TABLE agent_execution_logs (
    log_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    agent_name VARCHAR(100) NOT NULL,
    execution_type VARCHAR(50), -- SCHEDULED, USER_REQUEST, TRIGGERED
    trigger_source VARCHAR(100),
```

```
-- Execution Details
started_at TIMESTAMP NOT NULL,
completed_at TIMESTAMP,
duration_seconds INTEGER,
status VARCHAR(20), -- SUCCESS, FAILED, PARTIAL, TIMEOUT
```

```
-- Input/Output
input_params JSONB,
output_result JSONB,
error_message TEXT,
error_stack_trace TEXT,
```

```
-- Resource Usage
llm_calls_count INTEGER DEFAULT 0,
llm_tokens_used INTEGER DEFAULT 0,
llm_cost_usd DECIMAL(10,4),
db_queries_count INTEGER DEFAULT 0,
```

```
-- Metadata
environment VARCHAR(20), -- DEVELOPMENT, STAGING, PRODUCTION
```

```
version VARCHAR(50)
```

```
);
```

```
CREATE INDEX idx_agent_log_name_time ON agent_execution_logs(agent_name, started_at DESC);
```

```
CREATE INDEX idx_agent_log_status ON agent_execution_logs(status);
```

```
CREATE INDEX idx_agent_log_time ON agent_execution_logs(started_at DESC);
```

```
-- =====  
-- USER FEEDBACK & APPROVALS  
-- =====
```

```
CREATE TABLE user_approvals (
```

```
    approval_id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
```

```
    user_id VARCHAR(100) NOT NULL,
```

```
    user_name VARCHAR(255),
```

```
    -- What's being approved
```

```
    entity_type VARCHAR(50), -- PRICING_DECISION, WEEKLY_PLAN, CAMPAIGN
```

```
    entity_id UUID NOT NULL,
```

```
    -- Recommendation
```

```
    recommended_action TEXT,
```

```
    recommended_reasoning TEXT,
```

```
    confidence DECIMAL(4,3),
```

```
    -- User Decision
```

```
    decision VARCHAR(20), -- APPROVED, REJECTED, MODIFIED
```

```
    modified_values JSONB,
```

```
    user_reasoning TEXT,
```

```
    -- Timing
```

```
    requested_at TIMESTAMP NOT NULL,
```

```
    responded_at TIMESTAMP,
```

```
    response_time_seconds INTEGER,
```

```
    -- Learning
```

```
    feedback_quality_score DECIMAL(4,3),
```

```
    incorporated_into_learning BOOLEAN DEFAULT FALSE
```

```
);
```

```
CREATE INDEX idx_approval_user ON user_approvals(user_id, requested_at DESC);
```

```
CREATE INDEX idx_approval_entity ON user_approvals(entity_type, entity_id);
CREATE INDEX idx_approval_status ON user_approvals(decision);
```

```
-- =====
-- TRIGGERS FOR UPDATED_AT
-- =====
```

```
CREATE OR REPLACE FUNCTION update_updated_at_column()
RETURNS TRIGGER AS $$

BEGIN
    NEW.updated_at = CURRENT_TIMESTAMP;
    RETURN NEW;
END;

$$ language 'plpgsql';
```

```
CREATE TRIGGER update_sku_knowledge_updated_at BEFORE UPDATE ON sku_knowledge
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_pricing_episodes_updated_at BEFORE UPDATE ON pricing_episodes
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_weekly_plans_updated_at BEFORE UPDATE ON weekly_plans
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_campaign_episodes_updated_at BEFORE UPDATE ON campaign_episodes
    FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

---

## 5.2 Data Ingestion Pipeline

etl/ninjacart\_sync.py:



python

```
import pandas as pd
from sqlalchemy import create_engine
from datetime import datetime, timedelta
import logging
from typing import Dict, Optional
```

```
logger = logging.getLogger(__name__)
```

```
class NinjacartDataSync:
```

```
    """
```

ETL pipeline for syncing data from Ninjacart MySQL to PostgreSQL.

Purpose: Keep PostgreSQL data layer updated with latest operational data

Schedule: Runs daily at 6:00 AM (syncs previous day's data)

Process:

1. Connect to both databases
2. Extract data from Ninjacart MySQL
3. Transform to match PostgreSQL schema
4. Load into PostgreSQL (upsert)
5. Validate data quality

```
"""
```

```
def __init__(self, source_db_url: str, target_db_url: str):
    self.source_engine = create_engine(source_db_url)
    self.target_engine = create_engine(target_db_url)
    self.logger = logging.getLogger(__name__)
```

```
def sync_daily_actualls(self, date: datetime.date) -> Dict:
```

```
    """
```

Sync daily actualls from Ninjacart MySQL to PostgreSQL.

Input: date (datetime.date) - Date to sync

Output: Dict with sync stats

Process:

1. Query Ninjacart MySQL for all daily metrics
2. Calculate derived metrics (penetration, margins, etc.)
3. Upsert into PostgreSQL daily\_actualls table
4. Return sync statistics

.....

```
self.logger.info(f"Syncing daily actuals for {date}")

try:
    # Query Ninjacart MySQL
    query = f"""
SELECT
    DATE(so.created_at) as date,
    s.sku_id,
    s.sku_name,

    -- Sales metrics
    COUNT(DISTINCT so.customer_id) as customers_count,
    COUNT(DISTINCT so.order_id) as orders_count,
    SUM(soi.quantity_kg) as tpc,
    AVG(soi.quantity_kg) as tpo,

    -- Pricing
    AVG(soi.sale_price) as sp,
    AVG(pi.purchase_price) as pp,
    AVG(m.wholesale_price) as wsp,

    -- Inventory
    MAX(inv.current_stock_kg) as inventory_kg,

    -- Quality
    AVG(qc.rejection_rate) as qqc_rejection_pct,
    AVG(r.return_rate) as otpr_pct,

    -- External factors
    MAX(c.is_festival) as is_festival,
    MAX(w.temperature) as temperature,
    MAX(w.weather_condition) as weather_condition

FROM sale_orders so
JOIN sale_order_items soi ON so.order_id = soi.order_id
JOIN skus s ON soi.sku_id = s.sku_id
LEFT JOIN procurement_invoices pi ON s.sku_id = pi.sku_id
    AND DATE(pi.created_at) = '{date}'
LEFT JOIN market_prices m ON s.sku_id = m.sku_id
    AND DATE(m.created_at) = '{date}'
```

```

LEFT JOIN inventory inv ON s.sku_id = inv.sku_id
    AND DATE(inv.updated_at) = '{date}'
LEFT JOIN quality_checks qc ON s.sku_id = qc.sku_id
    AND DATE(qc.created_at) = '{date}'
LEFT JOIN returns r ON so.order_id = r.order_id
LEFT JOIN calendar c ON DATE(c.date) = '{date}'
LEFT JOIN weather w ON DATE(w.date) = '{date}'

WHERE DATE(so.created_at) = '{date}'
GROUP BY DATE(so.created_at), s.sku_id, s.sku_name
"""

df = pd.read_sql(query, self.source_engine)

if df.empty:
    self.logger.warning(f"No data found for {date}")
    return {"status": "no_data", "rows_synced": 0}

# Transform
df['ts'] = pd.to_datetime(df['date'])
df['is_weekend'] = df['ts'].dt.dayofweek >= 5

# Calculate penetration
total_active_customers = self._get_total_active_customers(date)
df['penetration'] = (df['customers_count'] / total_active_customers) * 100

# Calculate margins
df['pm'] = df['wsp'] - df['pp']
df['sm'] = df['sp'] - df['pp']
df['cm'] = df['wsp'] - df['wsp'] * 0.10 # 10% commission
df['gm'] = df['pm'] + df['sm'] - df['cm']

# Supply status
df['supply_status'] = df.apply(self._determine_supply_status, axis=1)
df['days_of_inventory'] = df['inventory_kg'] / df['tpc']

# Select columns for PostgreSQL
columns_to_sync = [
    'ts', 'sku_id', 'orders_count', 'customers_count', 'tpc', 'tpo',
    'penetration', 'wsp', 'pp', 'sp', 'pm', 'sm', 'cm', 'gm',
    'inventory_kg', 'supply_status', 'days_of_inventory',
]

```

```

'qqc_rejection_pct', 'otpr_pct', 'is_festival', 'is_weekend',
'weather_condition', 'temperature'
]

df_final = df[columns_to_sync]

# Upsert to PostgreSQL
rows_synced = self._upsert_to_postgres(df_final, 'daily_actuals')

self.logger.info(f"Successfully synced {rows_synced} rows for {date}")

return {
    "status": "success",
    "date": str(date),
    "rows_synced": rows_synced,
    "skus_processed": len(df_final['sku_id'].unique())
}

except Exception as e:
    self.logger.error(f"Error syncing data for {date}: {str(e)}")
    return {
        "status": "error",
        "date": str(date),
        "error": str(e)
    }

def _get_total_active_customers(self, date: datetime.date) -> int:
    """Calculate total active customers for penetration"""
    query = f"""
    SELECT COUNT(DISTINCT customer_id) as active_customers
    FROM sale_orders
    WHERE DATE(created_at) >= '{date - timedelta(days=30)}'
    AND DATE(created_at) <= '{date}'
    """
    result = pd.read_sql(query, self.source_engine)
    return int(result['active_customers'].iloc[0])

def _determine_supply_status(self, row) -> str:
    """Determine supply status based on inventory and demand"""
    # Simplified logic - enhance based on your business rules
    if pd.isna(row['inventory_kg']) or pd.isna(row['tpc']):

```

```

return 'UNKNOWN'

ratio = row['inventory_kg'] / row['tpc'] if row['tpc'] > 0 else 0

if ratio > 1.5:
    return 'SURPLUS'
elif ratio < 0.8:
    return 'DEFICIT'
else:
    return 'BALANCED'

def _upsert_to_postgres(self, df: pd.DataFrame, table_name: str) -> int:
    """Upsert data to PostgreSQL"""
    from sqlalchemy.dialects.postgresql import insert

    # Convert DataFrame to list of dicts
    records = df.to_dict('records')

    # Upsert
    rows_affected = 0
    with self.target_engine.begin() as conn:
        for record in records:
            stmt = insert(table_name).values(**record)
            stmt = stmt.on_conflict_do_update(
                index_elements=['ts', 'sku_id'],
                set_=record
            )
            result = conn.execute(stmt)
            rows_affected += result.rowcount

    return rows_affected

```

### Scheduled ETL Job (etl/scheduler.py):



python

```

from prefect import flow, task
from prefect.schedules import CronSchedule
from datetime import datetime, timedelta
import os

@task
def sync_yesterday_data():
    """Sync yesterday's data every morning at 6 AM"""
    yesterday = datetime.now().date() - timedelta(days=1)

    sync = NinjacartDataSync(
        source_db_url=os.getenv('NINJACART_DB_URL'),
        target_db_url=os.getenv('DATABASE_URL')
    )

    result = sync.sync_daily_actuals(yesterday)
    return result

@flow(name="daily_data_sync")
def daily_sync_flow():
    """Main flow for daily data synchronization"""
    result = sync_yesterday_data()

    if result['status'] == 'success':
        print(f"✅ Data sync completed: {result}")
    else:
        print(f"❌ Data sync failed: {result}")
        # TODO: Send alert to team

# Schedule: Every day at 6:00 AM IST
schedule = CronSchedule(cron="0 6 * * *", timezone="Asia/Kolkata")

if __name__ == "__main__":
    daily_sync_flow()

```

## 6. Memory Systems Implementation

### 6.1 Procedural Memory (Redis)

`memory/procedural_memory.py`:



python

```
import redis
import json
from typing import Dict, Any, Optional, List
import os
from datetime import datetime
```

## class ProceduralMemory:

"""

Stores and retrieves business rules, pricing workflows, configurations.

Storage: Redis (in-memory for fast access)

TTL: Rules persist forever, Weekly targets expire after 7 days

Key Patterns:

- rules:{sku\_id} → Pricing rules
- weekly\_targets:{sku\_id}:{week} → Weekly targets
- workflow:{workflow\_name} → Execution steps

"""

```
def __init__(self, redis_url: Optional[str] = None):
```

```
    self.redis_client = redis.from_url(
        redis_url or os.getenv('REDIS_URL'),
        decode_responses=True
    )
    self._ensure_rules_loaded()
```

```
def _ensure_rules_loaded(self):
```

```
    """Load initial rules from config file if not exists"""
    if not self.redis_client.exists('rules:TOMATO_LOCAL'):
        self._load_initial_rules()
```

```
def _load_initial_rules(self):
```

```
    """Load rules from config file"""
    config_path = os.path.join(os.path.dirname(__file__), './config/pricing_rules.json')
```

```
    with open(config_path, 'r') as f:
        rules = json.load(f)
        for sku_id, config in rules.items():
            self.set_rules(sku_id, config)
```

```
def get_rules(self, sku_id: str) -> Dict[str, Any]:
```

!!!!

Get pricing rules for a SKU.

Input: sku\_id (str)

Output: Dict with rules

Raises: ValueError if rules not found

!!!!

```
key = f"rules:{sku_id}"
```

```
rules_json = self.redis_client.get(key)
```

```
if not rules_json:
```

```
    raise ValueError(f"No rules found for SKU: {sku_id}")
```

```
return json.loads(rules_json)
```

```
def set_rules(self, sku_id: str, rules: Dict[str, Any]):
```

"""Store rules for a SKU (no expiration)"""

```
key = f"rules:{sku_id}"
```

```
self.redis_client.set(key, json.dumps(rules))
```

```
def get_weekly_targets(self, sku_id: str, week: str) -> Dict[str, Any]:
```

!!!!

Get weekly targets set by Weekly Planning Agent.

Input:

- sku\_id: str

- week: str (format: "2025-W42")

Output: Dict with weekly targets

Returns default targets if not found.

!!!!

```
key = f"weekly_targets:{sku_id}:{week}"
```

```
targets_json = self.redis_client.get(key)
```

```
if not targets_json:
```

*# Fallback to default from rules*

```
    return self._get_default_targets(sku_id)
```

```
return json.loads(targets_json)
```

```
def set_weekly_targets(
```

```
self,
sku_id: str,
week: str,
targets: Dict[str, Any]
):
```

Store weekly targets (expires after 7 days).

Called by Weekly Planning Agent after approval.

```
key = f"weekly_targets:{sku_id}:{week}"
self.redis_client.setex(
    key,
    604800, # 7 days in seconds
    json.dumps(targets)
)
```

```
def _get_default_targets(self, sku_id: str) -> Dict[str, Any]:
    """Fallback to default targets from rules"""
    rules = self.get_rules(sku_id)

    return {
        "daily_tpc_target_kg": 2400, # Default baseline
        "penetration_target_pct": rules['penetration_target']['daily_max'],
        "gm_target": rules['gm_target'],
        "tpo_target_kg": rules['tpo_target'],
        "strategy": "MAINTAIN"
    }
```

```
def get_workflow(self, workflow_name: str = "pricing") -> List[str]:
```

Get workflow execution steps.

Input: workflow\_name (str)

Output: List of step names

```
key = f"workflow:{workflow_name}"
workflow_json = self.redis_client.get(key)

if not workflow_json:
    # Return default workflow
```

```
return self._get_default_workflow(workflow_name)

return json.loads(workflow_json)

def _get_default_workflow(self, workflow_name: str) -> List[str]:
    """Default workflows"""
    workflows = {
        "pricing": [
            "fetch_weekly_context",
            "fetch_current_wsp",
            "run_demand_forecast",
            "check_inventory_levels",
            "check_returns_risk",
            "retrieve_similar_scenarios",
            "calculate_base_price",
            "apply_adjustments",
            "validate_constraints",
            "generate_reasoning",
            "calculate_confidence",
            "determine_approval_needed"
        ],
        "weekly_planning": [
            "gather_past_week_performance",
            "assess_supply_stability",
            "decide_strategy",
            "calculate_historical_baseline",
            "set_weekly_targets",
            "analyze_penetration_gap",
            "plan_campaigns",
            "generate_checkpoints",
            "create_plan_document"
        ],
        "campaign": [
            "identify_trigger",
            "analyze_customer_segments",
            "retrieve_similar_campaigns",
            "calculate_optimal_discount",
            "predict_outcomes",
            "generate_recommendation",
            "create_approval_request"
        ]
    }
```

```
}
```

```
return workflows.get(workflow_name, [])
```

```
def cache_session_data(
```

```
    self,
```

```
    session_id: str,
```

```
    data: Dict[str, Any],
```

```
    ttl_seconds: int = 3600
```

```
):
```

```
    """Cache session data (e.g., conversation context)"""
    key = f"session:{session_id}"
    self.redis_client.setex(key, ttl_seconds, json.dumps(data))
```

```
def get_session_data(self, session_id: str) -> Optional[Dict[str, Any]]:
```

```
    """Retrieve session data"""
    key = f"session:{session_id}"
    data_json = self.redis_client.get(key)
```

```
    if not data_json:
```

```
        return None
```

```
    return json.loads(data_json)
```

**Configuration File** (config/pricing\_rules.json):



```
json
```

```
{  
  "TOMATO_LOCAL": {  
    "penetration_target": {  
      "daily_min": 45,  
      "daily_max": 47,  
      "weekly": 60  
    },  
    "gm_target": [4, 5],  
    "tpo_target": 29,  
    "timing_rules": {  
      "price_hike_time": "07:00",  
      "price_drop_time": "09:30",  
      "review_time": "11:30"  
    },  
    "inventory_rules": {  
      "surplus_threshold": 0.20,  
      "deficit_threshold": -0.20,  
      "surplus_action": "REDUCE_SP_OR_CAMPAIGN",  
      "deficit_action": "INCREASE_SP",  
      "high_inventory_days": 2.0  
    },  
    "margin_rules": {  
      "min_sm_percentage": 10,  
      "min_gm_per_kg": 2,  
      "ideal_pp_formula": "WSP - (0.10 * WSP) - FML",  
      "commission_pct": 10,  
      "fml_typical": 1.50  
    },  
    "markup_factors": {  
      "high_inventory": 1.15,  
      "normal": 1.25,  
      "low_inventory": 1.35  
    },  
    "campaign_defaults": {  
      "lapsed_customer_days": [7, 21],  
      "conversion_rate_estimate": 0.18,  
      "avg_order_size_kg": 20,  
      "min_budget": 500,  
      "max_budget_pct_of_revenue": 0.05  
    }  
  },  
}
```

```
"POTATO_GRADED": {
    "penetration_target": {
        "daily_min": 31,
        "daily_max": 34,
        "weekly": 50
    },
    "gm_target": [4, 5],
    "tpo_target": 30,
    "timing_rules": {
        "price_hike_time": "07:00",
        "price_drop_time": "09:30",
        "review_time": "11:30"
    },
    "inventory_rules": {
        "surplus_threshold": 0.20,
        "deficit_threshold": -0.20,
        "surplus_action": "REDUCE_SP_OR_CAMPAIGN",
        "deficit_action": "INCREASE_SP",
        "high_inventory_days": 3.0
    },
    "margin_rules": {
        "min_sm_percentage": 10,
        "min_gm_per_kg": 2,
        "ideal_pp_formula": "WSP - (0.10 * WSP) - FML",
        "commission_pct": 10,
        "fml_typical": 1.20
    },
    "markup_factors": {
        "high_inventory": 1.15,
        "normal": 1.25,
        "low_inventory": 1.30
    }
},
"ONION_LOCAL": {
    "penetration_target": {
        "daily_min": 31,
        "daily_max": 34,
        "weekly": 48
    },
    "gm_target": [3, 4],
    "tpo_target": 26,
```

```

"timing_rules": {
    "price_hike_time": "07:00",
    "price_drop_time": "09:30",
    "review_time": "11:30"
},
"inventory_rules": {
    "surplus_threshold": 0.20,
    "deficit_threshold": -0.20,
    "surplus_action": "REDUCE_SP_OR_CAMPAIGN",
    "deficit_action": "INCREASE_SP",
    "high_inventory_days": 3.0
},
"margin_rules": {
    "min_sm_percentage": 8,
    "min_gm_per_kg": 1.5,
    "ideal_pp_formula": "WSP - (0.10 * WSP) - FML",
    "commission_pct": 10,
    "fml_typical": 1.20
},
"markup_factors": {
    "high_inventory": 1.12,
    "normal": 1.20,
    "low_inventory": 1.28
}
}
}
}

```

## 6.2 Vector Database (Pinecone)

**memory/vector\_store.py:**



python

```
import pinecone
from langchain.embeddings import OpenAIEMBEDDINGS
from langchain.vectorstores import Pinecone as LangchainPinecone
from typing import List, Dict, Optional
import os
import uuid
```

## class VectorMemory:

"""

Stores and retrieves historical scenarios using similarity search (RAG).

Storage: Pinecone (managed vector database)

Embedding: OpenAI text-embedding-ada-002 (1536 dimensions)

Use Cases:

- Find similar pricing scenarios
- Retrieve relevant campaign outcomes
- Historical pattern matching

"""

```
def __init__(self):
    pinecone.init(
        api_key=os.getenv('PINECONE_API_KEY'),
        environment=os.getenv('PINECONE_ENV')
    )

    self.index_name = os.getenv('PINECONE_INDEX', 'category-agent-episodes')
    self.embeddings = OpenAIEMBEDDINGS(
        openai_api_key=os.getenv('OPENAI_API_KEY'),
        model="text-embedding-ada-002"
    )

# Create index if doesn't exist
if self.index_name not in pinecone.list_indexes():
    pinecone.create_index(
        name=self.index_name,
        dimension=1536,
        metric="cosine",
        pods=1,
        pod_type="p1.x1"
    )
```

```
self.index = pinecone.Index(self.index_name)
self.vectorstore = LangchainPinecone(
    self.index,
    self.embeddings.embed_query,
    "text"
)
```

```
def store_pricing_episode(self, episode: Dict):
```

```
    """
```

Store a pricing episode for future RAG retrieval.

Input: episode (Dict) - Pricing episode data

Output: episode\_id (str)

Process:

1. Create searchable text representation
2. Generate embedding using OpenAI
3. Store in Pinecone with metadata

```
    """
```

# Create rich text description

```
text = self._create_episode_text(episode)
```

# Create metadata for filtering

```
metadata = {
    "episode_id": episode['episode_id'],
    "sku_id": episode['sku_id'],
    "date": episode['date'],
    "wsp": float(episode['context']['wsp']),
    "pp": float(episode['context']['pp']),
    "supply_status": float(episode['context']['supply_status']),
    "sp_recommended": float(episode['decision']['sp_recommended']),
    "sp_approved": float(episode['decision']['sp_approved']),
    "confidence": float(episode['decision']['confidence']),
    "tpc_actual": float(episode['outcome']['tpc_actual']),
    "tpc_predicted": float(episode['outcome']['tpc_predicted']),
    "accuracy": float(episode['outcome'].get('overall_accuracy', 0.0))
}
```

# Store

```
self.vectorstore.add_texts(
```

```
texts=[text],  
metadatas=[metadata],  
ids=[episode['episode_id']]  
)
```

```
return episode['episode_id']
```

```
def _create_episode_text(self, episode: Dict) -> str:
```

```
    """Create searchable text from episode"""
```

```
    ctx = episode['context']
```

```
    dec = episode['decision']
```

```
    out = episode['outcome']
```

```
    text = f"""
```

```
Pricing Episode for {episode['sku_id']} on {episode['date']}
```

Market Context:

- Wholesale Price (WSP): ₹{ctx['wsp']}/kg
- Purchase Price (PP): ₹{ctx['pp']}/kg
- Procurement Margin (PM): ₹{ctx['pm']}/kg
- Inventory: {ctx['inventory\_kg']}kg
- Supply Status: {ctx['supply\_status']}
- Inventory Surplus: {ctx['inventory\_surplus\_pct']}%

Pricing Decision:

- Recommended Sale Price (SP): ₹{dec['sp\_recommended']}/kg
- Approved Sale Price: ₹{dec['sp\_approved']}/kg
- Reasoning: {dec['sp\_reasoning']}
- Confidence: {dec['confidence']}
- Human Approved: {dec['human\_approved']}

Outcomes:

- Predicted TPC: {out['tpc\_predicted']}kg
- Actual TPC: {out['tpc\_actual']}kg
- Predicted Penetration: {out['penetration\_predicted']}%
- Actual Penetration: {out['penetration\_actual']}%
- Predicted GM: ₹{out['gm\_predicted']}/kg
- Actual GM: ₹{out['gm\_actual']}/kg
- Prediction Accuracy: {out.get('overall\_accuracy', 0)\*100}%

Key Learnings: {out.get('learning\_notes', 'N/A')}

!!!!

```
return text

def search_similar_pricing_scenarios(
    self,
    query: str,
    sku_id: Optional[str] = None,
    wsp_range: Optional[tuple] = None,
    supply_status: Optional[str] = None,
    k: int = 5,
    min_accuracy: float = 0.7
) -> List[Dict]:
```

!!!!

Search for similar historical pricing scenarios.

Input:

- query: str (natural language description)
- sku\_id: Optional[str] (filter by SKU)
- wsp\_range: Optional[tuple] (min\_wsp, max\_wsp)
- supply\_status: Optional[str] (SURPLUS/BALANCED/DEFICIT)
- k: int (number of results)
- min\_accuracy: float (only return episodes with accuracy >threshold)

Output: List of similar scenarios with scores

Example query:

"high inventory negative PM WSP around 35"

!!!!

# Build filter

filter\_dict = {}

if sku\_id:

filter\_dict["sku\_id"] = {"\$eq": sku\_id}

if wsp\_range:

```
    filter_dict["wsp"] = {
        "$gte": wsp_range[0],
        "$lte": wsp_range[1]
    }
```

```
if supply_status:
    filter_dict["supply_status"] = {"$eq": supply_status}

if min_accuracy:
    filter_dict["accuracy"] = {"$gte": min_accuracy}

# Search
results = self.vectorstore.similarity_search_with_score(
    query,
    k=k,
    filter=filter_dict if filter_dict else None
)

return [
    {
        "text": doc.page_content,
        "metadata": doc.metadata,
        "similarity_score": float(score)
    }
    for doc, score in results
]

def store_campaign_episode(self, campaign: Dict):
    """Store campaign episode for future retrieval"""
    text = self._create_campaign_text(campaign)

    metadata = {
        "campaign_id": campaign['campaign_id'],
        "sku_id": campaign['sku_id'],
        "campaign_type": campaign['campaign_type'],
        "segment_name": campaign['segment_name'],
        "conversion_rate_actual": float(campaign['conversion_rate_actual']),
        "roi_actual": float(campaign['actual_roi']),
        "success_score": float(campaign['overall_success_score'])
    }

    self.vectorstore.add_texts(
        texts=[text],
        metadatas=[metadata],
        ids=[campaign['campaign_id']]
    )
```

```
return campaign['campaign_id']

def _create_campaign_text(self, campaign: Dict) -> str:
    """Create searchable text from campaign"""
    text = f"""
        Campaign Episode for {campaign['sku_id']} on {campaign['campaign_date']}
```

Campaign Details:

- Type: {campaign['campaign\_type']}
- Trigger: {campaign['campaign\_trigger']}
- Segment: {campaign['segment\_name']} ({campaign['segment\_size']} customers)
- Duration: {campaign['duration\_hours']} hours

Predictions:

- Expected Conversion: {campaign['conversion\_rate\_predicted']\*100}%
- Expected Incremental TPC: {campaign['incremental\_tpc\_predicted']}kg
- Expected ROI: {campaign['expected\_roi']}x

Actuals:

- Actual Conversion: {campaign['conversion\_rate\_actual']\*100}%
- Actual Incremental TPC: {campaign['incremental\_tpc\_actual']}kg
- Actual ROI: {campaign['actual\_roi']}x
- Success Score: {campaign['overall\_success\_score']}

Key Learnings: {campaign.get('learning\_notes', 'N/A')}

"""

```
return text
```

```
def search_similar_campaigns(
    self,
    query: str,
    sku_id: Optional[str] = None,
    campaign_type: Optional[str] = None,
    min_success_score: float = 0.6,
    k: int = 5
) -> List[Dict]:
    """Search for similar successful campaigns"""
    filter_dict = {}
```

```

if sku_id:
    filter_dict["sku_id"] = {"$eq": sku_id}

if campaign_type:
    filter_dict["campaign_type"] = {"$eq": campaign_type}

if min_success_score:
    filter_dict["success_score"] = {"$gte": min_success_score}

results = self.vectorstore.similarity_search_with_score(
    query,
    k=k,
    filter=filter_dict if filter_dict else None
)

return [
    {
        "text": doc.page_content,
        "metadata": doc.metadata,
        "similarity_score": float(score)
    }
    for doc, score in results
]

```

**Due to the comprehensive nature of this document, I'll continue with the remaining sections (Agents 3-7, Orchestrator, API, Frontend, Testing, Deployment, Team Responsibilities) in the next part. The document is currently at ~15,000 words. Would you like me to:**

1. Continue with all remaining sections in one large artifact (will be ~40,000-50,000 words total)
2. Create multiple downloadable artifacts (Part 1: Foundation, Part 2: Agents, Part 3: Deployment)
3. Focus on specific sections you need most urgently

Let me know and I'll provide the complete guide!